

Toward Transparent and Parsimonious Methods for Automatic Performance Tuning

Pedro Bruel
phrb@ime.usp.br
July 9 2021

Introduction (5 min)

Trends on Hardware Design

- Hardware has ceased to provide "effortless" performance gains
- Performance continues to increase
- Accelerators are important, but suffer from the same scaling limits
- Code optimization is crucial for performance, and will continue to be

An Example of Autotuning: Loop Tiling and Unrolling

- How to restructure memory accesses in loops to increase throughput by leveraging cache locality?
- Size and shape of the resulting search space
- Introduce notation: $f : \mathcal{X} \rightarrow \mathbb{R}$

Autotuning Problems in Other Domains

- Number of parameters and combinations growing over time
- Earlier application to optimize BLAS routines
- Autotuning for specific domains and Neural Networks

Common Approaches to Autotuning

- Function minimization methods
 - Online learning?
- Surrogate-based optimization
 - Linear Models
 - Gaussian Processes
- Design of Experiments

Contributions of this Thesis

- Striving to develop and apply transparent and parsimonious autotuning methods
- Applications in this thesis:

Domain	Method
CUDA compiler parameters	Function minimization methods with Online Learning
FPGA compiler parameters	
OpenCL Laplacian Kernel	Function minimization methods, Linear Models, Gaussian Process Regression
SPAPT Kernels	Linear Models, Gaussian Process Regression
CNN Mixed-Precision Quantization	Gaussian Process Regression

Methods for Function Minimization (10 minutes)

Overview

- We know information about f :
- Derivative-based
- Other heuristics

Search Heuristics with Multi-Armed Bandit

- OpenTuner
- Ensemble of search heuristics
 - Coordinated by a MAB algorithm (online learning)

Application: High-Level Synthesis for FPGAs

Search Space and Performance Metrics

Results

Discussion

- Curse of dimensionality
- It is often unclear:
 - if there is something to find, and when to stop exploring
- It is impossible to:
 - interpret optimizations

Design of Experiments (15 minutes)

Linear Models

- Learning: Building surrogates, used for optimization
- Introduce notation:
 - $\hat{f}_\theta : \mathcal{X} \rightarrow \mathbb{R}$
 - Model of f : $f(\mathbf{x}) = \mathbf{x}^\top \theta + \varepsilon$
 - Surrogate $\hat{f}_\theta(\mathbf{x}) = \mathbf{x}^\top \hat{\theta}$.
- Best Linear Unbiased Estimator
- Learning methods assume the design \mathbf{X} is given

Design of Experiments

- Statistical methods to choose the design \mathbf{X} to minimize surrogate model variance
- Notation:
 - Factors, levels, design
- Simple linear model example for 2-factor designs
- Factorial designs , screening

Optimal Design

- Distributing points according to initial modeling hypotheses decreases model matrix determinant (associated with variance)
- Good for exploiting known search space structure, or verifying existing hypotheses

Space-filling Designs

- Curse of dimensionality for sampling:
 - Most sampled points will be on the "shell"
- LHS: Partition and then sample, need to optimize later
- Low-discrepancy: deterministic space-filling sequences

Interpreting Significance

- ANOVA for Linear Models
 - Isolate "significant" factors
- Sobol indices
 - expensive computation

A Transparent and Parsimonious Approach to Autotuning

- Explain paper diagram

Application: GPU Laplacian

Search Space and Performance Metrics

Results

- Comparison with multiple methods
- Leave GPR for later

Interpreting the Optimization

Application: SPAPT kernels

- Pick one?

Search Spaces and Performance Metrics

Results

- Is there anything to find?
- Leave GPR for later

Interpreting the Optimization

Discussion

- Motivating results in the Laplacian kernel
- It is possible to interpret results, guide optimization
 - sometimes simpler models give better results
- For SPAPT kernels, it is still unclear:
 - if there is something to find, and when to stop exploring
 - is there a global optimum, is it "hidden"?
 - how to find it, if so? (can learning do it?)
- Random Sampling has good performance
 - Abundance of local optima?
- What is the most effective level of abstraction for optimizing a program?
 - Compiler, kernel, machine, model, dependencies?

Gaussian Process Regression (10 minutes)

More Flexibility with Gaussian Process Regression

- Introduce notation:
 - Model of f : $f(\mathbf{x}) \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$
 - Surrogate $\hat{f}_{\theta}(\mathbf{x}) \sim f(\mathbf{x}) \mid \mathbf{X}, \mathbf{y}$

Expected Improvement: Balancing Exploitation and Exploration

- How to decide where to measure next?

Application: GPU Laplacian and SPAPT

- GPR was applied to these problems too

Results: GPU Laplacian

- GPR is good too, but the simpler model is more consistent

Results: SPAPT

- GPR still can't find better configurations

Application: Quantization for Convolutional Neural Networks

Search Space, Constraints, and Performance Metrics

- Comparing with a Reinforcement Learning approach in the original paper
- ImageNet

Results

Interpreting the Optimization

- Sobol indices, inconclusive

Discussion

- Low-discrepancy sampling in high dimension
- Constraints complicate exploration
- Multi-objective optimization
- A more complex method usually produces less interpretable results, but not always achieves better optimizations

Conclusion (5 min)



Contributions of this Thesis

- Striving to develop and apply transparent and parsimonious autotuning methods
- Applications in this thesis:

Domain	Method
CUDA compiler parameters	Function minimization methods with Online Learning
FPGA compiler parameters	
OpenCL Laplacian Kernel	Function minimization methods, Linear Models, Gaussian Process Regression
SPAPT Kernels	Linear Models, Gaussian Process Regression
CNN Mixed-Precision Quantization	Gaussian Process Regression

Reproducibility of Performance Tuning Experiments

- Redoing all the work for different problems
- Complementary approaches:
 - Completely evaluate small sets of a search space
 - Collaborative optimizing for different architectures, problems

Key Discussions

Curse of Dimensionality for Autotuning Problems

- Implications for Sampling and Learning
- Space-filling helps, but does not solve
- Constraints

Which method to use?

- Design of Experiments for transparency and parsimony when building and interpreting statistical models
- Linear models for simpler spaces and problems
- Gaussian Process Surrogates for more complex situations

It is often unclear if there is something to find

- Abundance of local optima
- Is there a global optimum, is it "hidden"?
 - How to find it, if so? (can learning do it?)
- What is the most effective level of abstraction for optimizing a program?
 - Compiler, kernel, machine, model, dependencies?
- When to stop?

Conclusion

Toward Transparent and Parsimonious Methods for Automatic Performance Tuning

Pedro Bruel
phrb@ime.usp.br
July 9 2021