

Towards Transparent and Parsimonious Methods for Automatic Performance Tuning

Pedro Bruel

October 24, 2020

Contents

I	Introduction	5
1	The Need for Autotuning	7
1.1	Historical Trends in Hardware Design	9
1.2	Loop Nest Optimization as an Autotuning Problem	12
1.3	Characterizing Search Spaces	16
1.4	Thesis Contributions	19
2	Efforts for Reproducible Science	23
2.1	Introduction	23
2.2	Computational Documents with Emacs and Org mode	23
2.3	Versioning for Code, Text, and Data	23
II	Autotuning using Function Minimization Methods	25
3	Methods for Known or Cheap-to-Evaluate Functions	29
3.1	Stochastic Methods	29
3.2	Methods Based on Derivatives	29
4	Methods for Unknown or Expensive-to-Evaluate Functions	31
4.1	Statistical Learning Methods	31
4.2	Design of Experiments	32
4.3	Gaussian Process Regression	32
5	Applications to Autotuning	33

5.1	Search Heuristics	33
5.2	Statistical Learning	33
5.3	Design of Experiments	33
5.4	Gaussian Process Regression	33
III	Experimental Validation	35
6	Search Heuristics	37
6.1	Results and Limits of Analysis	37
7	A Design of Experiments Methodology	39
7.1	Inference for Autotuning under a Tight Budget	39
7.2	Prediction and Inference for Source-to-Source Transformation	39
8	Gaussian Process Regression: A more Flexible Method	41
8.1	Expressing Structure with Covariance Kernels?	41
8.2	Results with Quantization of Deep Neural Network Layers	41
9	Conclusion	43

Part I

Introduction

The Need for Autotuning

High Performance Computing has been a cornerstone of scientific and industrial progress for at least five decades. By paying the cost of increased complexity, software and hardware engineering advances continue to overcome several challenges on the way of the sustained performance improvements observed during the last fifty years. A consequence of this mounting complexity is that reaching the theoretical peak hardware performance for a given program requires not only expert knowledge of specific hardware architectures, but also mastery of programming models and languages for parallel and distributed computing.

If we state performance optimization problems as *search* or *learning* problems, by converting implementation and configuration choices to *parameters* which might affect performance, we can draw from and adapt proven methods from search, mathematical optimization, and statistical learning. The effectiveness of these adapted methods on performance optimization problems varies greatly, and hinges on practical and mathematical properties of the problem and the corresponding *search space*. The application of such methods to the automation of performance tuning for specific hardware, under a set of *constraints*, is named *autotuning*.

Improving performance also relies on gathering application-specific knowledge, which entails extensive experimental costs since, with the exception of linear algebra routines, theoretical peak performance is not always a reachable comparison baseline. When adapting methods for autotuning we must face challenges emerging from practical properties, such as restricted time and cost budgets, constraints on feasible parameter values, and the need to mix *categorical*, *continuous*, and *discrete* parameters. To achieve useful results we must also choose methods that make hypotheses compatible with problem search spaces, such as the existence of *discoverable*, or at least *exploitable*, relationships between parameters and performance. Choosing an autotuning method requires balancing the exploration of a problem, that is, seeking to discover and explain relationships between parameters and performance, and the exploitation of known or discovered relationships, seeking only to find the best possible performance.

Search algorithms based on machine learning heuristics are not the best candidates for autotuning domains where measurements are lengthy and costly, such as compiling industrial-level FPGA programs, because these algorithms rely on the availability of a large number of measurements. They also assume good optimizations are reachable from a starting position, and that tendencies observed locally in the search space are exploitable. These assumptions are not usually true in common autotuning domains, as shown in the work of Seymour *et al.* [1].

Autotuning search spaces also usually have non-linear constraints and undefined regions, which are also expected to decrease the effectiveness of search based on heuristics and machine learning. An additional downside to heuristics- and machine learning-based search is that, usually, optimization choices cannot be explained, and knowledge gained during optimization is not reusable. The main contribution of this thesis is to study how to overcome the reliance on these assumptions about search spaces, and the lack of explainable optimizations, from the point of view of a *Design of Experiments* (DoE), or *Experimental Design*, methodology to autotuning.

One of the first detailed descriptions and mathematical treatment of DoE was presented by Ronald Fisher [2] in his 1937 book *The Design of Experiments*, where he discussed principles of experimentation, latin square sampling and factorial designs. Later books such as the ones from Jain [3], Montgomery [4] and Box *et al.* [5] present comprehensive and detailed foundations. Techniques based on DoE are *parsimonious* because they allow decreasing the number of measurements required to determine certain relationships between parameters and metrics, and are *transparent* because parameter selections and configurations can be justified by the results of statistical tests.

In DoE terminology, a *design* is a plan for executing a series of measurements, or *experiments*, whose objective is to identify relationships between *factors* and *responses*. While factors and responses can refer to different concrete entities in other domains, in computer experiments factors can be configuration parameters for algorithms and compilers, for example, and responses can be the execution time or memory consumption of a program.

Designs can serve diverse purposes, from identifying the most significant factors for performance, to fitting analytical performance models for the response. The field of DoE encompasses the mathematical formalization of the construction of experimental designs. More practical works in the field present algorithms to generate designs with different objectives and restrictions.

The contributions of this thesis are strategies to apply to program autotuning the DoE methodology, and *Gaussian Process Regression* [6]. This thesis presents background and a high-level view of the theoretical foundations of each method, and detailed discussions of the challenges involved in specializing the general definitions of search heuristics and statistical learning methods to different autotuning problems, as well as what can be *learned*




Figure 1.1: 49 years of microprocessor data, highlighting the sustained exponential increases and reductions on transistor counts and fabrication processes, the stagnation of frequency scaling around 2005, and one solution found for it, the simultaneous exponential increase on logical core count. Data from Wikipedia [9, 10]

about specific autotuning search spaces, and how that acquired knowledge can be leveraged for further optimization.

This chapter aims to substantiate the claim that autotuning methods have a fundamental role to play on the future of program performance optimization, arguing that the value and the difficulty of the efforts to carefully tune software became more apparent ever since advances in hardware stopped leading to effortless performance improvements, at least from the programmer’s perspective. The following sections discuss the historical context for the changes in trends on computer architecture, and characterize the search spaces found when optimizing performance on different domains.

1.1 Historical Trends in Hardware Design

The physical constraints imposed by technological advances on circuit design were evident since the first vacuum tube computers that already spanned entire floors, such as the ENIAC in 1945 [7]. The practical and economical need to fit more computing power into real estate is one force for innovation in hardware design that spans its history, and is echoed in modern supercomputers, such as the *Summit* from *Oak Ridge National Laboratory* [8], which spans an entire room.

Figure 1.1 highlights the unrelenting and so far successful pursuit of smaller transistor fabrication processes, and the resulting capability to fit more computing power on a fixed chip area. This trend was already observed in integrated circuits by Gordon Moore *et al.*

in 1965 [11], who also postulated its continuity. The performance improvements produced by the design efforts to make Moore's forecast a self-fulfilling prophecy were boosted until around 2005 by the performance gained from increases in circuit frequency.

Robert Dennard *et al.* remarked in 1974 [12] that smaller transistors, in part because they generate shorter circuit delays, decrease the energy required to power a circuit and enable an increase in operation frequency without breaking power usage constraints. This scaling effect, named *Dennard's scaling*, is hindered primarily by leakage current, caused by quantum tunneling effects in small transistors. Figure 1.1 shows a marked stagnation on frequency increase after around 2005, as transistors crossed the 10^2nm fabrication process. It was expected that leakage due to tunneling would limit frequency scaling strongly, even before the transistor fabrication process reached 10nm [13].

Current hardware is now past the effects of Dennard's scaling. The increase in logical cores around 2015 can be interpreted as preparation for and mitigation of the end of frequency scaling, and ushered in an age of multicore scaling. Still, in order to meet power consumption constraints, up to half of a multicore processor could have to be powered down, at all times. This phenomenon is named *Dark Silicon* [14], and presents significant challenges to current hardware designers and programmers [15, 16, 17].

The *Top500* [18] list gathers information about commercially available supercomputers, and ranks them by performance on the LINPACK benchmark [19]. Figure 1.2 shows the peak theoretical performance R_{Peak} , and the maximum performance achieved on the LINPACK benchmark R_{Max} , in $Tflops/s$, for the top-ranked supercomputers on TOP500. Despite the smaller performance gains from hardware design that are to be expected for post-Dennard's scaling processors, the increase in computer performance has sustained an exponential climb, sustained mostly by software improvements.

Although *hardware accelerators* such as GPUs and FPGAs, have also helped to support exponential performance increases, their use is not an escape from the fundamental scaling constraints imposed by current semiconductor design. Figure 1.3 shows the increase in processor and accelerator core count on the top-ranked supercomputers on Top500. Half of the top-ranked supercomputers in the last decade had accelerator cores and, of those, all had around ten times more accelerator than processor cores. The apparent stagnation of core count in top-ranked supercomputers, even considering accelerators, highlights the crucial impact software optimization has on performance.

Advances in hardware design are currently not capable of providing performance improvements via frequency scaling without dissipating more power than the processor was designed to support, which violates power constraints and risks damaging the circuit. From the programmer's perspective, effortless performance improvements from hardware have not been expected for quite some time, and the key to sustaining historical trends in performance scaling has lied in accelerators, parallel and distributed programming




Figure 1.2: Sustained exponential increase of theoretical *RPeak* and achieved *RMax* performance for the supercomputer ranked 1st on TOP500 [18]




Figure 1.3: Processor and accelerator core count in supercomputers ranked 1st on TOP500 [18]. Core count trends for supercomputers are not necessarily bound to processor trends observed on Figure 1.1.

libraries, and fine tuning of several stages of the software stack, from instruction selection to the layout of neural networks.

The problem of optimizing software for performance presents its own challenges. The search spaces that emerge from autotuning problems grow quickly to a size for which it would take a prohibitive amount of time to determine the best configuration by exhaustively evaluating all possibilities. Although this means we must seek to decrease the amount of possibilities, by restricting allowed parameter values, or dropping parameters completely, it

is often unclear how to decide which parameters should be restricted or dropped. The next sections introduce a simple autotuning problem, present an overview of the magnitude of the dimension of autotuning search spaces, and briefly introduce the methods commonly used to explore search spaces, some of which are discussed in detail in Chapter II.

1.2 Loop Nest Optimization as an Autotuning Problem

Algorithms for linear algebra problems are fundamental to scientific computing and statistics. Therefore, decreasing the execution time of algorithms such as general matrix multiplication (GEMM) [20], and others from the original BLAS [21], is an interesting and well motivated example, that we will use to introduce the autotuning problem.

One way to improve the performance of such linear algebra programs is to exploit cache locality by reordering and organizing loop iterations, using source code transformation methods such as loop *tiling*, or *blocking*, and *unrolling*. We will now briefly describe loop tiling and unrolling for a simple linear algebra problem. After, we will discuss an autotuning search space for blocking and unrolling applied to GEMM, and how these transformations generate a relatively large and complex search space, which we can explore using autotuning methods.

Figure 1.4 shows three versions of code in the C language that, given three square matrices A , B , and C , computes $C = C + A + B^T$. The first optimization we can make is to preemptively load to cache, or *prefetch*, as many as possible of the elements we know will be needed at any given iteration, as is shown in Figure 1.4a. The shaded elements on the top row of Figure 1.5 represent the elements that could be prefetched in iterations of Figure 1.4a.

Since C matrices are stored in *row-major order*, each access of an element of B forces loading the next row elements, even if we explicitly prefetch a column of B . Since we are accessing B in a *column-major order*, the prefetched row elements would not be used until we reached the corresponding column. Therefore, the next column elements will have to be loaded at each iteration, considerably slowing down the computation.

We can solve this problem by reordering memory accesses to request only prefetched elements. It suffices to adequately split loop indices into blocks, as shown in Figure 1.4b. Now, memory accesses are performed in *tiles*, as shown on the bottom row of Figure 1.5. If blocks are correctly sized to fit in cache, we can improve performance by explicitly prefetching each tile. After blocking, we can still improve performance by *unrolling* loop iterations, which forces register usage and helps the compiler to identify regions that can be *vectorized*. A conceptual implementation of loop unrolling is shown in Figure 1.4c.

Looking at the loop nest optimization problem from the autotuning perspective, the two *parameters* that emerge from the implementations are the *block size*, which controls the

```

int N = 256;
float A[N][N], B[N][N], C[N][N];
int i, j;
// Initialize A, B, C
for(i = 0; i < N; i++){
    // Load line i of A to fast memory
    for(j = 0; j < N; j++){
        // Load C[i][j] to fast memory
        // Load column j of B to fast memory
        C[i][j] += A[i][j] + B[j][i];
        // Write C[i][j] to main memory
    }
}

int N = 256;
int B_size = 4;
int A[N][N], B[N][N], C[N][N];
int i, j, x, y;
// Initialize A, B, C
for(i = 0; i < N; i += B_size){
    for(j = 0; j < N; j += B_size){
        // Load block (i, j) of C to fast memory
        // Load block (i, j) of A to fast memory
        // Load block (j, i) of B to fast memory
        for(x = i; x < min(i + B_size, N); x++){
            for(y = j; y < min(j + B_size, N); y++){
                C[x][y] += A[x][y] + B[y][x];
            }
        }
        // Write block (i, j) of C to main memory
    }
}

```

(a) Regular implementation

(b) Blocked, or tiled

```

int N = 256;
int B_size = 4;
int A[N][N], B[N][N], C[N][N];
int i, j, k;
// Initialize A, B, C
for(i = 0; i < N; i += B_size){
    for(j = 0; j < N; j += B_size){
        // Load block (i, j) of C to fast memory
        // Load block (i, j) of A to fast memory
        // Load block (j, i) of B to fast memory
        C[i + 0][j + 0] += A[i + 0][j] * B[i][j + 0];
        C[i + 0][j + 1] += A[i + 0][j] * B[i][j + 1];
        // Unroll the remaining 12 iterations
        C[i + B_size - 1][j + B_size - 2] += A[i + B_size - 1][j] * B[i][j + B_size - 2];
        C[i + B_size - 1][j + B_size - 1] += A[i + B_size - 1][j] * B[i][j + B_size - 1];
        // Write block (i, j) of C to main memory
    }
}

```

(c) Tiled and unrolled

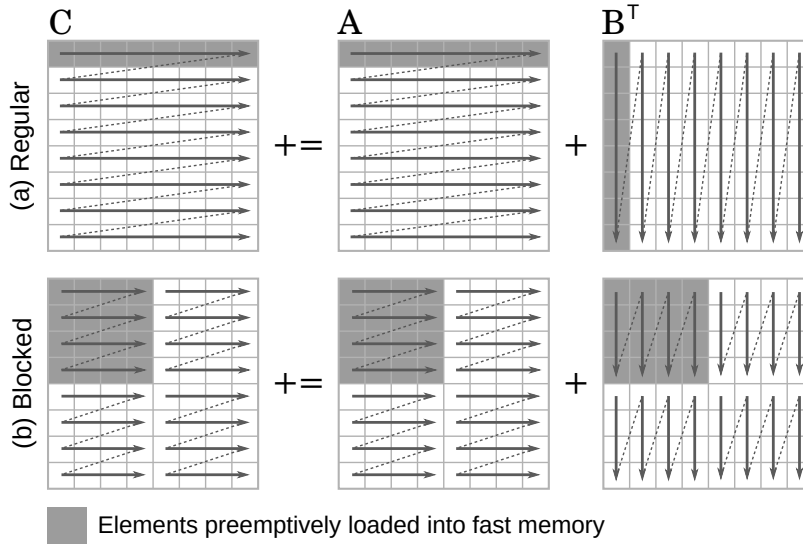
Figure 1.4: Loop nest optimizations for $C = C + A + B^T$, in C

Figure 1.5: Access patterns for matrices in $C = C + A + B^T$, with loop nest optimizations. Panel (a) shows the access order of a regular implementation, and panel (b) shows the effect of loop tiling, or blocking

stride, and the *unrolling factor*, which controls the number of unrolled iterations. Larger block sizes are desirable, because we want to avoid extra comparisons, but blocks should

```

int N = 256;
float A[N][N], B[N][N], C[N][N];
int i, j, k;
// Initialize A, B, C
for(i = 0; i < N; i++){
    // Load line i of A to fast memory
    for(j = 0; j < N; j++){
        // Load C[i][j] to fast memory
        // Load column j of B to fast memory
        for(k = 0; k < N; k++){
            C[i][j] += A[i][k] * B[k][j];
        }
        // Write C[i][j] to main memory
    }
}

```

(a) Regular implementation

```

int N = 256;
int B_size = 4;
int A[N][N], B[N][N], C[N][N];
int i, j, k, x, y;
// Initialize A, B, C
for(i = 0; i < N; i += B_size){
    for(j = 0; j < N; j += B_size){
        // Load block (i, j) of C to fast memory
        for(k = 0; k < N; k++){
            // Load block (i, k) of A to fast memory
            // Load block (k, j) of B to fast memory
            for(x = i; x < min(i + B_size, N); x++){
                for(y = j; y < min(j + B_size, N); y++){
                    C[x][y] += A[x][k] * B[k][y];
                }
            }
        }
        // Write block (i, j) of C to main memory
    }
}

```

(b) Blocked, or tiled

```

int N = 256;
int B_size = 4;
int A[N][N], B[N][N], C[N][N];
int i, j, k;
// Initialize A, B, C
for(i = 0; i < N; i += B_size){
    for(j = 0; j < N; j += B_size){
        // Load block (i, j) of C to fast memory
        for(k = 0; k < N; k++){
            // Load block (i, k) of A to fast memory
            // Load block (k, j) of B to fast memory
            C[i + 0][j + 0] += A[i + 0][k] * B[k][j + 0];
            C[i + 0][j + 1] += A[i + 0][k] * B[k][j + 1];
            // Unroll the remaining 12 iterations
            C[i + B_size - 1][j + B_size - 2] += A[i + B_size - 1][k] * B[k][j + B_size - 2];
            C[i + B_size - 1][j + B_size - 1] += A[i + B_size - 1][k] * B[k][j + B_size - 1];
        }
        // Write block (i, j) of C to main memory
    }
}

```

(c) Tiled and unrolled

Figure 1.6: Loop nest optimizations for GEMM, in C

be small enough to ensure access to as few as possible out-of-cache elements. Likewise, the unrolling factor should be large, to leverage vectorization and available registers, but not so large that it forces memory to the stack.

The values of block size and unrolling factor that optimize performance will depend on the cache hierarchy, register layout, and vectorization capabilities of the target processor, but also on the memory access pattern of the target algorithm. In addition to finding the best values for each parameter independently, an autotuner must ideally aim to account for the *interactions* between parameters, that is, for the fact that the best value for each parameter might also depend on the value chosen for the other.

The next loop optimization example comes from Seymour *et al.* [1], and considers 128 blocking and unrolling values, in the interval [0, 127], for the GEMM algorithm. The three panels of Figure 1.6 show conceptual implementations of loop blocking and unrolling for GEMM in C. A block size of zero results in the implementation from Figure 1.6a, and an unrolling factor of zero performs a single iteration per condition check.

It is straightforward to change the block size of the implementations from Figure 1.6, but the unrolling factor is not exposed as a parameter. To test different unrolling values




Figure 1.7: An exhaustively measured search space, defined by loop blocking and unrolling parameters, for a sequential GEMM kernel. Reproduced from Seymour *et al.* [1]

we need to generate new versions of the source code with different numbers of unrolled iterations. We can do that with code generators or with *source-to-source transformation* tools [22, 23, 24]. It is often necessary to modify the program we wish to optimize in order to provide a configuration interface and expose its implicit parameters. Once we are able to control the block size and the loop unrolling factor, we determine the target search space by choosing the values to be explored.

In this example, the search space is defined by the $128^2 = 16384$ possible combinations of blocking and unrolling values. The performance of each combination in the search space, shown in *Mflops/s* in Figure 1.7, was measured for a sequential GEMM implementation, using square matrices of size 400 [1]. We can represent this autotuning search space as a *3D landscape*, since we have two configurable parameters and a single target performance metric. In this setting, the objective is to find the *highest* point, since the objective is to *maximize* *Mflops/s*, although usually the performance metric is transformed so that the objective is its *minimization*.

On a first look, there seems to be no apparent global search space structure in the landscape on Figure 1.7, but local features jump to the eyes, such as the “valley” across all block sizes for low unrolling factors, the “ramp” across all unrolling factors for low block sizes, and the series of jagged “plateaus” across the middle regions, with ridges for identical or divisible block sizes and unrolling factors. A careful look reveals also that there is a curvature along the unrolling factor axis. Also of note is the abundance in this landscape of *local minima*, that is, points with relatively good performance, surrounded by points with worse performance. By exhaustively evaluating all possibilities, the original study determined that the best performance on this program was achieved with a block size of 80 and an unrolling factor of 2.

In this conceptual example, all $\approx 1.64 \times 10^4$ configurations were exhaustively evaluated, but it is impossible to do so in most settings where autotuning methods are useful. The

next section provides a perspective of the autotuning domains and methods employed in current research, presenting a selection of search spaces and discussing the trends that can be observed on search space size, targeted HPC domains, and chosen optimization methods.

1.3 Characterizing Search Spaces

Autotuning methods have been used to improve performance in an increasingly large variety of domains, from the earlier applications to linear algebra subprograms, to the now ubiquitous construction and configuration of neural networks, to the configuration of the increasingly relevant tools for the re-configurable hardware of FPGAs. In this setting, it is not far-fetched to establish a link between the continued increases in performance and hardware complexity, that we discussed previously in this chapter, to the increases in dimension and size of the autotuning problems that we can now tackle.

Figure 1.8 presents search space dimension, measured as the number of parameters involved, and size, measured as the number of possible parameter combinations, for a selection of search spaces from 14 autotuning domains. Precise information about search space characterization is often missing from works on autotuning methods and applications. The characterization of most of the search spaces in Figure 1.8 was obtained directly from the text of the corresponding published paper, but for some it was necessary to extract characterizations from the available source code. Still, it was impossible to obtain detailed descriptions of search spaces for many of the published works on autotuning methods and applications, and in that way the sample shown in this section is biased, because it contains only information on works that provided it.

The left hand panel of Figure 1.8 shows search spaces with up to 60 parameters. The over-representation of search spaces for linear algebra domains in this sample stands out on the left hand panel, but the domain is not present on the remaining portion of the sample, shown on the right hand panel. The largest search spaces for which we were able to find information on published work are defined for the domains of neural network configuration, High-Level Synthesis for FPGAs, compiler parameters, and domain-specific languages.

None of the largest search spaces in this sample, that is, the ones outside the zoomed area of the left hand panel, come from works earlier than 2009. The sustained performance improvements we discussed previously have enabled and pushed autotuning research toward progressively larger problems, which has also been done to most research areas. Increased computing power has made it feasible, or at least tolerable, to apply search heuristics and statistical learning methods to find program configurations that improve performance.

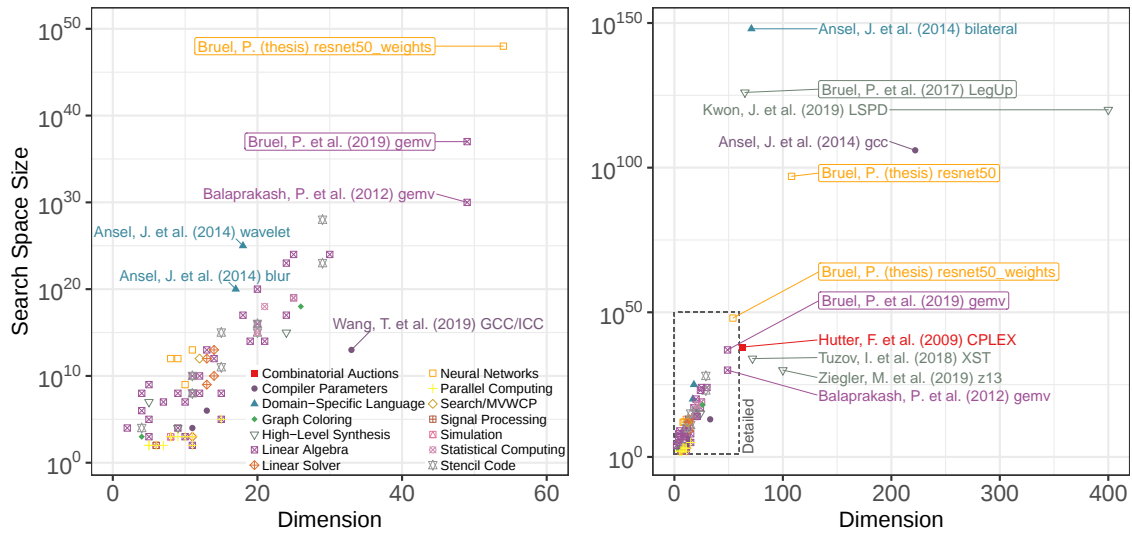


Figure 1.8: Dimension and search space size for autotuning problems from 14 domains [25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 1] The left panel shows a zoomed view of the right panel

It is straightforward to produce an extremely large autotuning search space. Compilers have hundreds of binary flags that can be considered for selection, generating a large set of combinations. Despite that, regarding performance improvements, it is likely that most configuration parameters will have a small impact, that is, that only a handful of parameters are responsible for changes in performance. Search spaces are often much more restrictive than the one we discussed in Section 1.2. Autotuning problem definitions usually come with *constraints* on parameter values and limited *experimental budgets*, and *runtime* failures for some configurations are often unpredictable. In this context, finding configurations that improve performance and determining the subset of *significant parameters* are considerable challenges.

Search heuristics, such as methods based on genetic algorithms and gradient descent, are a natural way to tackle these challenges because they consist of procedures for exploiting existing and unknown relationships between parameters and performance without making or requiring *explicit hypotheses* about the problem. Despite that, most commonly used heuristics make *implicit hypotheses* about search spaces which are not always verified, such as assuming that good configurations are *reachable* from a random starting point.

Autotuning methods that make explicit hypotheses about the target program, such as methods based on Design of Experiments, require some initial knowledge, or willingness to make assumptions, about underlying relationships, and are harder to adapt to constrained scenarios, but have the potential to produce *explainable* optimizations. In general, methods based on Machine Learning have enough flexibility to perform well in complex search spaces and make few assumptions about problems, but usually provide little, if any, that

Table 1.1: Autotuning methods used by a sample of systems, in different domains, ordered by publishing year. Methods were classified as either Search Heuristics (SH), Machine Learning (ML), or more precisely when the originating work provided detailed information. Earlier work favored employing Search Heuristics, which are less prominent in recent work, which favors methods based on Machine Learning.

System	Domain	Method	Year
PhiPAC [45]	Linear Algebra	SH (Exhaustive)	1997
ATLAS [46]	Linear Algebra	SH (Exhaustive)	1998
FFTW [47]	Digital Signal Processing	SH (Exhaustive)	1998
Active Harmony [48]	Domain-Specific Language	SH	2002
OSKI [49]	Linear Algebra	SH	2005
Seymour, K. <i>et al.</i> [1]	Linear Algebra	SH	2008
PRO [36]	Linear Algebra	SH	2009
ParamILS [37]	Combinatorial Auctions	SH	2009
PetaBricks [24]	Domain-Specific Language	SH (Genetic Algorithm)	2009
MILEPOST GCC [50]	Compiler Parameters	ML	2011
Orio [25]	Linear Algebra	ML (Decision Trees)	2012
pOSKI [27]	Linear Algebra	SH	2012
INSIEME [51]	Compiler Parameters	SH (Genetic Algorithm)	2012
OpenTuner [26]	Compiler Parameters	SH	2014
Lgen [52]	Linear Algebra	SH	2014
OPAL [53]	Parallel Computing	SH	2014
Mametjanov, A. <i>et al.</i> [33]	High-Level Synthesis	ML (Decision Trees)	2015
CLTune [54]	Parallel Computing	SH	2015
Guerreiro, J. <i>et al.</i> [55]	Parallel Computing	SH	2015
Collective Mind [56]	Compiler Parameters	ML	2015
Abdelfattah, A. <i>et al.</i> [34]	Linear Algebra	SH (Exhaustive)	2016
TANGRAM [57]	Domain-Specific Language	SH	2016
MASE-BDI [58]	Environmental Land Change	SH	2016
Xu, C. <i>et al.</i> [35]	High-Level Synthesis	SH	2017
Apollo [59]	Parallel Computing	ML (Decision Trees)	2017
DeepHyper [29]	Neural Networks	ML (Decision Trees)	2018
Tuzov, I. <i>et al.</i> [39]	High-Level Synthesis	Design of Experiments	2018
Periscope [41]	Compiler Parameters	SH	2018
SynTunSys [40]	High-Level Synthesis	SH	2019
Kwon, J. <i>et al.</i> [42]	High-Level Synthesis	ML	2019
FuncyTuner [43]	Compiler Parameters	SH	2019
Ol'ha, J. <i>et al.</i> [44]	Parallel Computing	Sensitivity Analysis	2019
Petrovic, F. <i>et al.</i> [28]	Linear Algebra	SH	2020
Chu, Y. <i>et al.</i> [38]	Search/MVWCP	SH	2020

can be used to explain optimization choices or derive relationships between parameters and performance.

Table 1.1 lists some autotuning systems, their target domains, and the employed method, ordered by publication date. Some systems that did not provide detailed search space descriptions and could not be included in Figure 1.8, especially some of the earlier work, provided enough information to categorize their autotuning methods. In contrast,

many more recent works, especially those using methods based on Machine Learning, did not provide specific method information. Earlier work often deals with search spaces small enough to exhaustively evaluate, and using search heuristics to optimize linear algebra programs is the most prominent category of earlier work in this sample. Later autotuning work target more varied domains, with the most prominent domains in this sample being parallel computing, compiler parameters, and High-Level Synthesis. Systems using methods based on Machine Learning become more common on later work than systems using heuristics.

Chapter 5 provides more detailed definitions and discussions of the applicability, effectiveness, and explanatory power of autotuning methods based on search heuristics and statistical learning. The remainder of this chapter details the contributions of this thesis and the structure of this document.

1.4 Thesis Contributions

Computer performance has sustained exponential increases over the last half century, despite current physical limits on hardware design. Software optimization has performed an increasingly substantial role on performance improvement over the last 15 years, requiring the exploration of larger and more complex search spaces than ever before.

Autotuning methods are one approach to tackle performance optimization of complex search spaces, enabling exploitation of existing relationships between program parameters and performance. We can derive autotuning methods from well-established statistics, although their usage is not common in or standardized for autotuning domains.

Initial work on this thesis studied the effectiveness of classical and standard search heuristics, such as Simulated Annealing, on autotuning problems. The first target autotuning domain was the set of parameters of a compiler for CUDA programs. The search heuristics for this case study were implemented using the OpenTuner framework [26], and consisted of an ensemble of search heuristics coordinated by a Multi-Armed Bandit algorithm. The autotuner searched for a set of compilation parameters that optimized 17 heterogeneous GPU kernels, from a set of approximately 10^{23} possible combinations of all parameters. With 1.5h autotuning runs we have achieved up to 4× speedup in comparison with the CUDA compiler’s high-level optimizations. The compilation and execution times of programs in this autotuning domain are relatively fast, and were in the order of a few seconds to a minute. Since measurement costs are relatively small, search heuristics could find good optimizations using as many measurements as needed. A detailed description of this work is available in our paper [32] published in the *Concurrency and Computation: Practice and Experience* journal.

The next case study was developed in collaboration with *Hewlett-Packard Enterprise*, and consisted of applying the same heuristics-based autotuning approach to the configuration

of parameters involved in the generation of FPGA hardware specification from source code in the C language, a process called *High-Level Synthesis* (HLS). The main difference from our work with GPU compiler parameters was the time to obtain the hardware specification, which could be in the order of hours for a single kernel.

In this more complex scenario, we achieved up to 2× improvements for different hardware metrics using conventional search algorithms. These results were obtained in a simple HLS benchmark, for which compilation times were in the order of minutes. The search space was composed of approximately 10^{123} possible configurations, which is much larger than the search space in our previous work with GPUs. Search space size and the larger measurement cost meant that we did not expect the heuristics-based approach to have the same effectiveness as in the GPU compiler case study. This work was published [60] at the 2017 *IEEE International Conference on ReConFigurable Computing and FPGAs*.

Approaches using classical machine learning and optimization techniques would not scale to industrial-level HLS, where each compilation can take hours to complete. Search space properties also increase the complexity of the problem, in particular its structure composed of binary, factorial and continuous variables with potentially complex interactions. Our results on autotuning HLS for FPGAs corroborate the conclusion that the empirical autotuning of expensive-to-evaluate functions, such as those that appear on the autotuning of HLS, require a more parsimonious and transparent approach, that can potentially be achieved using the DoE methodology. The next section describes our work on applying the DoE methodology to autotuning.

The main contribution of this thesis is a strategy to apply the Design of Experiments methodology to autotuning problems. The strategy is based on linear regression and its extensions, Analysis of Variance (ANOVA), and Optimal Design. The strategy requires the formulation of initial assumptions about the target autotuning problem, which are refined with data collected by efficiently selected experiments. The main objectives are to identify relationships between parameters, suggesting regions for further experimentation in a transparent way, that is, in a way that is supported by statistical tests of significance. The effectiveness of the proposed strategy, and its ability to explain optimizations it finds, are evaluated on autotuning problems in the source-to-source transformation domain. The initial stages of this work resulted in a publication on IEEE/ACM CCGrid [30].

Further efforts in this direction were dedicated to describe precisely what can be learned about search spaces from the application of the methodology, and to refine the differentiation of the approach for the sometimes conflicting objectives of model assessment and prediction. These discussions are presented in Chapter 7.

Because the Design of Experiments methodology requires the specification of a class of initial performance models, the methodology can sometimes achieve worse prediction capabilities when there is considerable uncertainty on initial assumptions about the underlying relationships. Chapter 8 describes the application to autotuning of Gaussian Process

Regression, an approach that trades some explanatory power for a much larger and more flexible class of underlying models. We evaluate the performance of this approach on the source-to-source transformation problem from Chapter 7, and on larger autotuning problem on the quantization of *Deep Neural Network* (DNN) layers.

Efforts for Reproducible Science

2.1 Introduction

2.2 Computational Documents with Emacs and Org mode

2.3 Versioning for Code, Text, and Data

Part II

Autotuning using Function Minimization Methods

The effectiveness of search heuristics on autotuning can be limited [1, 61, 62], between other factors, by underlying hypotheses about the search space, such as the reachability of the global optimum and the smoothness of search space surfaces, which are frequently not respected. The derivation of relationships between parameters and performance from search heuristic optimizations is greatly hindered, if not rendered impossible, by the biased way these methods explore parameters. Some parametric learning methods, such as Design of Experiments, are not widely applied to autotuning. These methods perform structured parameter exploration, and can be used to build and validate performance models, generating transparent and cost-effective optimizations [33, 30]. Other methods from the parametric family are more widely used, such as Bandit Algorithms [35]. Nonparametric learning methods, such as Decision Trees [63] and Gaussian Process Regression [64], are able to reduce model bias greatly, at the expense of increased prediction variance. Figure ?? categorizes some autotuning methods according to some of the key hypotheses and branching questions underlying each method.

Methods for Known or Cheap-to-Evaluate Functions

3.1 Stochastic Methods

1. Single-State Methods
2. Methods Based on Populations

3.2 Methods Based on Derivatives

1. Local Descent Methods

Methods for Unknown or Expensive-to-Evaluate Functions

4.1 Statistical Learning Methods

1. Introduction

- (a) Optimization with Surrogate Models
- (b) The Bias-Variance Trade-Off
- (c) Inference and prediction
- (d) Supervised and Unsupervised Learning
- (e) Parametric and Nonparametric Statistics

2. Supervised Methods for Regression and Classification

- KNN, Linear Regression, Logistic Regression, Neural Networks

3. Model Assessment

- Train and Test Sets
- Train and Test Mean Squared Error
- Cross Validation, Bootstrapping
- Information Criteria for Model Assessment
 - Mallow's C_p and AIC
 - BIC

4. Choosing Experiments X : An Unposed Question

5. Linear Regression

- Ridge, Lasso
- (a) Ordinary Least Squares: Gauss-Markov (BLUE)

(b) Transformations of X

6. Classification of Statistical Learning Methods

4.2 Design of Experiments

1. Introduction

(a) Posing the Question of Choosing Experiments X

i. Reducing Experimental Cost

ii. Improving Model Inference and Prediction

2. Screening

3. Factorial Designs

4. Space-Filling Designs

5. Optimal Design

(a) Mixing Categorical and Numerical Factors

(b) Optimality Criteria

(c) Exchange Algorithms for Optimal Design Construction

6. Inference with Analysis of Variance (ANOVA)

7. Sampling under Search Space Constraints

8. Classification of Design of Experiments Methods

4.3 Gaussian Process Regression

1. Introduction

2. Sampling Functions from Multidimensional Gaussian Distributions

3. Nonparametric Modeling with Covariance Kernels

4. Sensitivity Analysis with Sobol Indices

5. Other Nonparametric Methods

6. Classification of Nonparametric Methods

Applications to Autotuning

5.1 Search Heuristics

1. OpenTuner: Leveraging Ensembles of Heuristics
 - (a) The Multi-Armed Bandit, Area Under the Curve Approach

5.2 Statistical Learning

1. Software for Autotuning with Statistical Learning
 - (a) The R Language
2. Results using Statistical Learning
 - (a) Linear Regression and Design of Experiments
 - (b) Methods based on Machine Learning

5.3 Design of Experiments

1. Inference for Autotuning under a Tight Budget
 - (a) Fractional Factorial Designs for Chip Design on FPGAs

5.4 Gaussian Process Regression

Part III

Experimental Validation

Search Heuristics

6.1 Results and Limits of Analysis

1. GPU Compiler Parameters
2. High-Level Synthesis for FPGAs
3. Limitations of Analyzing Results from Search Heuristics

A Design of Experiments Methodology

7.1 Inference for Autotuning under a Tight Budget

1. Screening for Main Effects of GPU Compiler Parameters
2. Optimal Designs for a GPU Laplacian Kernel
3. Space-Filling Designs for Exploring Quantization of DNN Layers

7.2 Prediction and Inference for Source-to-Source Transformation

1. Results with SPAPT Kernels

Gaussian Process Regression: A more Flexible Method

8.1 Expressing Structure with Covariance Kernels?

8.2 Results with Quantization of Deep Neural Network Layers

Conclusion

Bibliography

- [1] K. Seymour, H. You, and J. Dongarra, "A comparison of search heuristics for empirical code optimization." in *CLUSTER*, 2008, pp. 421–429.
- [2] R. A. Fisher, *The design of experiments*. Oliver And Boyd; Edinburgh; London, 1937.
- [3] P. N. D. Bukh, *The art of computer systems performance analysis, techniques for experimental design, measurement, simulation and modeling*. JSTOR, 1992.
- [4] D. C. Montgomery, *Design and analysis of experiments*. John wiley & sons, 2017.
- [5] G. E. Box, J. S. Hunter, and W. G. Hunter, *Statistics for experimenters: design, innovation, and discovery*. Wiley-Interscience New York, 2005, vol. 2.
- [6] C. K. Williams and C. E. Rasmussen, *Gaussian processes for machine learning*. MIT press Cambridge, MA, 2006, vol. 2, no. 3.
- [7] P. E. Ceruzzi, E. Paul *et al.*, *A history of modern computing*. MIT press, 2003.
- [8] O. R. N. Laboratory, "Summit – Oak Ridge Leadership Computing Facility," <https://www.olcf.ornl.gov/olcf-resources/compute-systems/summit>, 2020, [Online; accessed 08-Jun-2020].
- [9] Wikipedia, "Transistor count," https://en.wikipedia.org/wiki/Transistor_count, 2020, [Online; accessed 08-Jun-2020].
- [10] —, "Microprocessor Chronology," https://en.wikipedia.org/wiki/Microprocessor_chronology, 2020, [Online; accessed 08-Jun-2020].
- [11] G. E. Moore *et al.*, "Cramming more components onto integrated circuits," 1965.
- [12] R. H. Dennard, F. H. Gaensslen, H.-N. Yu, V. L. Rideout, E. Bassous, and A. R. LeBlanc, "Design of ion-implanted mosfet's with very small physical dimensions," *IEEE Journal of solid-state circuits*, vol. 9, no. 5, pp. 256–268, 1974.
- [13] D. J. Frank, R. H. Dennard, E. Nowak, P. M. Solomon, Y. Taur, and H.-S. P. Wong, "Device scaling limits of si mosfets and their application dependencies," *Proceedings of the IEEE*, vol. 89, no. 3, pp. 259–288, 2001.

- [14] H. Esmaeilzadeh, E. Blem, R. S. Amant, K. Sankaralingam, and D. Burger, "Dark silicon and the end of multicore scaling," in *2011 38th Annual international symposium on computer architecture (ISCA)*. IEEE, 2011, pp. 365–376.
- [15] S. Venkataramani, S. T. Chakradhar, K. Roy, and A. Raghunathan, "Approximate computing and the quest for computing efficiency," in *2015 52nd ACM/EDAC/IEEE Design Automation Conference (DAC)*. IEEE, 2015, pp. 1–6.
- [16] H.-Y. Cheng, J. Zhan, J. Zhao, Y. Xie, J. Sampson, and M. J. Irwin, "Core vs. uncore: The heart of darkness," in *2015 52nd ACM/EDAC/IEEE Design Automation Conference (DAC)*. IEEE, 2015, pp. 1–6.
- [17] J. Henkel, H. Khdr, S. Pagani, and M. Shafique, "New trends in dark silicon," in *2015 52nd ACM/EDAC/IEEE Design Automation Conference (Dac)*. IEEE, 2015, pp. 1–6.
- [18] Top500, "Top500 List," <https://www.top500.org/>, 2020, [Online; accessed 08-Jun-2020].
- [19] J. J. Dongarra, P. Luszczek, and A. Petitet, "The LINPACK benchmark: past, present and future," *Concurrency and Computation: practice and experience*, vol. 15, no. 9, pp. 803–820, 2003.
- [20] J. J. Dongarra, J. Du Croz, S. Hammarling, and I. S. Duff, "A set of level 3 basic linear algebra subprograms," *ACM Transactions on Mathematical Software (TOMS)*, vol. 16, no. 1, pp. 1–17, 1990.
- [21] C. L. Lawson, R. J. Hanson, D. R. Kincaid, and F. T. Krogh, "Basic linear algebra subprograms for fortran usage," *ACM Transactions on Mathematical Software (TOMS)*, vol. 5, no. 3, pp. 308–323, 1979.
- [22] B. Videau, K. Pouget, L. Genovese, T. Deutsch, D. Komatitsch, F. Desprez, and J.-F. Méhaut, "BOAST: A metaprogramming framework to produce portable and efficient computing kernels for hpc applications," *The International Journal of High Performance Computing Applications*, p. 1094342017718068, 2017.
- [23] A. Hartono, B. Norris, and P. Sadayappan, "Annotation-based empirical performance tuning using Orio," in *Parallel & Distributed Processing, 2009. IPDPS 2009. IEEE International Symposium on*. IEEE, 2009, pp. 1–11.
- [24] J. Ansel, C. Chan, Y. L. Wong, M. Olszewski, Q. Zhao, A. Edelman, and S. Amarasinghe, "Petabricks: a language and compiler for algorithmic choice," *ACM Sigplan Notices*, vol. 44, no. 6, pp. 38–49, 2009.
- [25] P. Balaprakash, S. M. Wild, and B. Norris, "SPAPT: Search problems in automatic performance tuning," *Procedia Computer Science*, vol. 9, pp. 1959–1968, 2012.

- [26] J. Ansel, S. Kamil, K. Veeramachaneni, J. Ragan-Kelley, J. Bosboom, U.-M. O'Reilly, and S. Amarasinghe, "Opentuner: An extensible framework for program autotuning," in *Proceedings of the 23rd international conference on Parallel architectures and compilation*. ACM, 2014, pp. 303–316.
- [27] J.-H. Byun, R. Lin, K. A. Yelick, and J. Demmel, "Autotuning sparse matrix-vector multiplication for multicore," *EECS, UC Berkeley, Tech. Rep*, 2012.
- [28] F. Petrovič, D. Štřelák, J. Hozzová, J. Ol'ha, R. Trembecký, S. Benkner, and J. Filipovič, "A benchmark set of highly-efficient cuda and opencl kernels and its dynamic autotuning with kernel tuning toolkit," *Future Generation Computer Systems*, 2020.
- [29] P. Balaprakash, M. Salim, T. Uram, V. Vishwanath, and S. Wild, "Deephyper: Asynchronous hyperparameter search for deep neural networks," in *2018 IEEE 25th International Conference on High Performance Computing (HiPC)*. IEEE, 2018, pp. 42–51.
- [30] P. Bruel, S. Quinito Masnada, B. Videau, A. Legrand, J.-M. Vincent, and A. Goldman, "Autotuning under Tight Budget Constraints: A Transparent Design of Experiments Approach," in *The 19th Annual IEEE/ACM International Symposium in Cluster, Cloud, and Grid Computing (CCGrid 2019)*. IEEE/ACM, 2019.
- [31] P. Bruel, M. Amarís, and A. Goldman, "Autotuning gpu compiler parameters using opentuner," in *Proceedings of the WSCAD (Simpósio em Sistemas Computacionais de Alto Desempenho)*, 2015.
- [32] —, "Autotuning CUDA compiler parameters for heterogeneous applications using the OpenTuner framework," *Concurrency and Computation: Practice and Experience*, pp. e3973–n/a, 2017.
- [33] A. Mametjanov, P. Balaprakash, C. Choudary, P. D. Hovland, S. M. Wild, and G. Sabin, "Autotuning fpga design parameters for performance and power," in *Field-Programmable Custom Computing Machines (FCCM), 2015 IEEE 23rd Annual International Symposium on*. IEEE, 2015, pp. 84–91.
- [34] A. Abdelfattah, A. Haidar, S. Tomov, and J. Dongarra, "Performance, design, and autotuning of batched gemm for gpus," in *International Conference on High Performance Computing*. Springer, 2016, pp. 21–38.
- [35] C. Xu, G. Liu, R. Zhao, S. Yang, G. Luo, and Z. Zhang, "A parallel bandit-based approach for autotuning fpga compilation," in *Proceedings of the 2017 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays*. ACM, 2017, pp. 157–166.
- [36] A. Tiwari, C. Chen, J. Chame, M. Hall, and J. K. Hollingsworth, "A scalable auto-tuning framework for compiler optimization," in *Parallel & Distributed Processing, 2009. IPDPS 2009. IEEE International Symposium on*. IEEE, 2009, pp. 1–12.

- [37] F. Hutter, H. H. Hoos, K. Leyton-Brown, and T. Stützle, "ParamILS: an automatic algorithm configuration framework," *Journal of Artificial Intelligence Research*, vol. 36, no. 1, pp. 267–306, 2009.
- [38] Y. Chu, C. Luo, H. H. Hoos, Q. Lin, and H. You, "Improving the performance of stochastic local search for maximum vertex weight clique problem using programming by optimization," *arXiv*, pp. arXiv-2002, 2020.
- [39] I. Tuzov, D. de Andrés, and J.-C. Ruiz, "Tuning synthesis flags to optimize implementation goals: Performance and robustness of the leon3 processor as a case study," *Journal of Parallel and Distributed Computing*, vol. 112, pp. 84–96, 2018.
- [40] M. M. Ziegler, H.-Y. Liu, G. Gristede, B. Owens, R. Nigaglioni, J. Kwon, and L. P. Carloni, "Syntunsys: A synthesis parameter autotuning system for optimizing high-performance processors," in *Machine Learning in VLSI Computer-Aided Design*. Springer, 2019, pp. 539–570.
- [41] M. Gerndt, S. Benkner, E. César, C. Navarrete, E. Bajrovic, J. Dokulil, C. Guillén, R. Mijakovic, and A. Sikora, "A multi-aspect online tuning framework for hpc applications," *Software Quality Journal*, vol. 26, no. 3, pp. 1063–1096, 2018.
- [42] J. Kwon, M. M. Ziegler, and L. P. Carloni, "A learning-based recommender system for autotuning design flows of industrial high-performance processors," in *2019 56th ACM/IEEE Design Automation Conference (DAC)*. IEEE, 2019, pp. 1–6.
- [43] T. Wang, N. Jain, D. Beckingsale, D. Boehme, F. Mueller, and T. Gamblin, "Funcytuner: Auto-tuning scientific applications with per-loop compilation," in *Proceedings of the 48th International Conference on Parallel Processing*, 2019, pp. 1–10.
- [44] J. Ol'ha, J. Hozzová, J. Fousek, and J. Filipovič, "Exploiting historical data: pruning autotuning spaces and estimating the number of tuning steps," in *European Conference on Parallel Processing*. Springer, 2019, pp. 295–307.
- [45] J. Bilmes, K. Asanovic, C.-W. Chin, and J. Demmel, "Optimizing matrix multiply using PHiPAC: a portable, high-performance, ansi c coding methodology," in *Proceedings of International Conference on Supercomputing, Vienna, Austria*, 1997.
- [46] J. J. Dongarra and C. R. Whaley, "Automatically tuned linear algebra software (ATLAS)," *Proceedings of SC*, vol. 98, 1998.
- [47] M. Frigo and S. G. Johnson, "FFTW: An adaptive software architecture for the fft," in *Acoustics, Speech and Signal Processing, 1998. Proceedings of the 1998 IEEE International Conference on*, vol. 3. IEEE, 1998, pp. 1381–1384.
- [48] C. Țăpuș, I.-H. Chung, J. K. Hollingsworth *et al.*, "Active harmony: Towards automated performance tuning," in *Proceedings of the 2002 ACM/IEEE conference on Supercomputing*. IEEE Computer Society Press, 2002, pp. 1–11.

- [49] R. Vuduc, J. W. Demmel, and K. A. Yelick, "OSKI: A library of automatically tuned sparse matrix kernels," in *Journal of Physics: Conference Series*, vol. 16, no. 1. IOP Publishing, 2005, p. 521.
- [50] G. Fursin, Y. Kashnikov, A. W. Memon, Z. Chamski, O. Temam, M. Namolaru, E. Yom-Tov, B. Mendelson, A. Zaks, E. Courtois *et al.*, "Milepost gcc: Machine learning enabled self-tuning compiler," *International journal of parallel programming*, vol. 39, no. 3, pp. 296–327, 2011.
- [51] H. Jordan, P. Thoman, J. J. Durillo, S. Pellegrini, P. Gschwandtner, T. Fahringer, and H. Moritsch, "A multi-objective auto-tuning framework for parallel codes," in *High Performance Computing, Networking, Storage and Analysis (SC), 2012 International Conference for*. IEEE, 2012, pp. 1–12.
- [52] D. G. Spampinato and M. Püschel, "A basic linear algebra compiler," in *Proceedings of Annual IEEE/ACM International Symposium on Code Generation and Optimization*. ACM, 2014, p. 23.
- [53] C. Audet, K.-C. Dang, and D. Orban, "Optimization of algorithms with opal," *Mathematical Programming Computation*, vol. 6, no. 3, pp. 233–254, 2014.
- [54] C. Nugteren and V. Codreanu, "Clitune: A generic auto-tuner for opencl kernels," in *Embedded Multicore/Many-core Systems-on-Chip (MCSoc), 2015 IEEE 9th International Symposium on*. IEEE, 2015, pp. 195–202.
- [55] J. Guerreiro, A. Ilic, N. Roma, and P. Tomás, "Multi-kernel auto-tuning on gpus: Performance and energy-aware optimization," in *Parallel, Distributed and Network-Based Processing (PDP), 2015 23rd Euromicro International Conference on*. IEEE, 2015, pp. 438–445.
- [56] G. Fursin, A. Memon, C. Guillon, and A. Lokhmotov, "Collective mind, part ii: Towards performance-and cost-aware software engineering as a natural science," *arXiv preprint arXiv:1506.06256*, 2015.
- [57] L.-W. Chang, I. El Hajj, C. Rodrigues, J. Gómez-Luna, and W.-m. Hwu, "Efficient kernel synthesis for performance portable programming," in *Microarchitecture (MICRO), 2016 49th Annual IEEE/ACM International Symposium on*. IEEE, 2016.
- [58] C. G. Coelho, C. G. Abreu, R. M. Ramos, A. H. Mendes, G. Teodoro, and C. G. Ralha, "Mase-bdi: agent-based simulator for environmental land change with efficient and parallel auto-tuning," *Applied Intelligence*, vol. 45, no. 3, pp. 904–922, 2016.
- [59] D. Beckingsale, O. Pearce, I. Laguna, and T. Gamblin, "Apollo: Reusable models for fast, dynamic tuning of input-dependent code," in *The 31th IEEE International Parallel and Distributed Processing Symposium*, 2017.

- [60] P. Bruel, A. Goldman, S. R. Chalamalasetti, and D. Milojicic, "Autotuning High-Level Synthesis for FPGAs Using OpenTuner and LegUp," in *International Conference on Reconfigurable Computing and FPGAs (ReConFig)*, 2017.
- [61] P. Balaprakash, S. M. Wild, and P. D. Hovland, "Can search algorithms save large-scale automatic performance tuning?" in *ICCS*, 2011, pp. 2136–2145.
- [62] —, "An experimental study of global and local search algorithms in empirical performance tuning," in *International Conference on High Performance Computing for Computational Science*. Springer, 2012, pp. 261–269.
- [63] P. Balaprakash, A. Tiwari, S. M. Wild, and P. D. Hovland, "AutoMOMML: Automatic Multi-objective Modeling with Machine Learning," in *High Performance Computing: 31st International Conference, ISC High Performance 2016, Frankfurt, Germany, June 19-23, 2016, Proceedings*, M. J. Kunkel, P. Balaji, and J. Dongarra, Eds. Springer International Publishing, 2016, pp. 219–239.
- [64] M. Parsa, A. Ankit, A. Ziabari, and K. Roy, "PABO: Pseudo Agent-Based Multi-Objective Bayesian Hyperparameter Optimization for Efficient Neural Accelerator Design," in *2019 IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*. IEEE, 2019, pp. 1–8.