

# Autotuning Methods

Pedro Buel

June 11, 2020



# Contents

<b>1</b>	<b>The Need for Autotuning in High Performance Computing</b>	<b>5</b>
1.1	Historical Trends in Hardware Design . . . . .	6
1.2	Characterizing Search Spaces from Autotuning Problems . . . . .	10
1.3	Thesis Contributions . . . . .	10
1.4	Text Structure . . . . .	10
<b>2</b>	<b>Reproducible Science</b>	<b>13</b>
<b>3</b>	<b>Search and Statistical Learning for Autotuning</b>	<b>15</b>
<b>4</b>	<b>Design of Experiments</b>	<b>17</b>
<b>5</b>	<b>Gaussian Process Regression</b>	<b>19</b>
5.1	References . . . . .	19



# The Need for Autotuning in High Performance Computing

High Performance Computing has been a cornerstone of scientific and industrial progress for at least five decades. By paying the cost of increased complexity, software and hardware engineering advances continue to overcome several challenges on the way of the sustained performance improvements observed during the last fifty years. A consequence of this mounting complexity is that reaching the theoretical peak hardware performance for a given program requires not only expert knowledge of specific hardware architectures, but also mastery of programming models and languages for parallel and distributed computing.

If we state performance optimization problems as *search* or *learning* problems, by converting implementation and configuration choices to *parameters* which might affect performance, we can draw and adapt proven methods from search, mathematical optimization, and *statistical learning*. The effectiveness of these adapted methods on performance optimization problems varies greatly, and hinges on practical and mathematical properties of the problem and the corresponding *search space*. The application of such methods to the automation of program optimization for specific hardware, under a set of *constraints*, is named *autotuning*.

Improving performance also relies on gathering application-specific knowledge, which entails extensive experimental costs since, with the exception of linear algebra routines, theoretical peak performance is not always a reachable comparison baseline. When adapting methods for autotuning we must face challenges emerging from practical properties, such as restricted time and cost budgets, constraints on feasible parameter values, and the need to mix *categorical*, *continuous*, and *discrete* parameters. To achieve

useful results we must also choose methods that make hypotheses compatible with problem search spaces, such as the existence of *discoverable*, or at least *exploitable*, relationships between parameters and performance. Choosing an autotuning method requires determining a balance between the exploration of a problem, when we would seek to discover and explain relationships between parameters and performance, and the exploitation of the best optimizations we can find, when we would seek only to minimize performance.

The contributions of this thesis are strategies to apply to program autotuning the statistical learning methods of *Design of Experiments* [1], or *Experimental Design*, and *Gaussian Process Regression* [2]. This thesis presents background and a high-level view of the theoretical foundations of each method, and detailed discussions of the challenges involved in specializing the general definitions of search heuristics and statistical learning methods to different autotuning problems, as well as what can be *learned* about specific autotuning search spaces, and how that acquired knowledge can be leveraged for further optimization.

This chapter aims to substantiate the claim that autotuning methods have a fundamental role to play on the future of program performance optimization, arguing that the value and the difficulty of the efforts to carefully tune software became more apparent ever since advances in hardware stopped leading to effortless performance improvements, at least from the programmer’s perspective. The following sections discuss the historical context for the changes in trends on computer architecture, and characterize the search spaces found when optimizing performance on different domains.

## 1.1 Historical Trends in Hardware Design

The physical constraints imposed by technological advances on circuit design were evident since the first vacuum tube computers that already spanned entire floors, such as the ENIAC in 1945 [3]. The practical and economical need to fit more computing power into real estate is one force for innovation in hardware design that spans its history, and is echoed in modern supercomputers, such as the *Summit* from *Oak Ridge National Laboratory* [4], which spans an entire room.

Figure 1.1 highlights the unrelenting and so far successful pursuit of smaller transistor fabrication processes, and the resulting capability to fit more computing power on a fixed chip area. This trend was already observed in integrated circuits by Gor-

don Moore *et al.* in 1965 [5], who also postulated its continuity. The performance improvements produced by the design efforts to make Moore's forecast a self-fulfilling prophecy were boosted until around 2005 by the performance gained from increases in circuit frequency.




Figure 1.1: 49 years of microprocessor data, highlighting the sustained exponential increases and reductions on transistor counts and fabrication processes, the stagnation of frequency scaling around 2005, and one solution found for it, the simultaneous exponential increase on logical core count. Data from Wikipedia [6, 7]

Robert Dennard *et al.* remarked in 1974 [8] that smaller transistors, in part because they generate shorter circuit delays, decrease the energy required to power a circuit and enable an increase in operation frequency without breaking power usage constraints. This scaling effect, named *Dennard's scaling*, is hindered primarily by leakage current, caused by quantum tunneling effects in small transistors. Figure 1.1 shows a marked stagnation on frequency increase after around 2005, as transistors crossed the  $10^2\text{nm}$  fabrication process. It was expected that leakage due to tunneling would limit frequency scaling strongly, even before the transistor fabrication process reached 10nm [9].

Current hardware is now past the effects of Dennard's scaling. The increase in logical cores, still around 2015, can be interpreted as preparation for and mitigation of the end of frequency scaling, and ushered in an age of multicore scaling. Still, in order to meet power consumption constraints, up to half of a multicore processor could have to be powered down at all times. This phenomenon is named *Dark Silicon* [10], and still

presents significant challenges to hardware designers and programmers [11, 12, 13].

The *Top500* [14] list gathers information about commercially available supercomputers, and ranks them by their performance on the *LINPACK* benchmark [15]. Figure 1.2 shows the peak theoretical performance  $RPeak$ , and the maximum performance achieved on the LINPACK benchmark  $RMax$ , in  $Tflops/s$ , for the top-ranked supercomputers on TOP500. Despite the smaller performance gains from hardware design that are to be expected for post-Dennard's scaling processors, the increase in computer performance has continued its exponential climb, sustained mostly by software improvements.




Figure 1.2: Sustained exponential increase of theoretical  $RPeak$  and achieved  $RMax$  performance for the supercomputer ranked 1<sup>st</sup> on TOP500 [14]

Although hardware accelerators, such as GPUs and FPGAs, have also helped to sustain exponential performance increase, their use is not an escape from the fundamental scaling constraints imposed by current semiconductor design. Figure 1.3 shows the increase in processor and accelerator core count on the top-ranked supercomputers on Top500. Half of the top-ranked supercomputers in the last decade had accelerator cores, and of those, all had around ten times more accelerator than processor cores. The apparent stagnation of core count in top-ranked supercomputers, even considering accelerators, highlights the crucial impact of software on performance.

Advances in hardware design are currently not capable of providing performance improvements via frequency scaling without dissipating more power than the proces-



sor was designed to support, which violates power constraints and risks damaging the circuit. From the programmer’s perspective, effortless performance improvements from hardware have not been expected for quite some time, and the key to sustaining historical trends in performance scaling has lied in accelerators, parallel and distributed programming libraries, and fine tuning of several stages of the software stack, from instruction selection to the layout of neural networks.




Figure 1.3: Processor and accelerator core count in supercomputers ranked 1<sup>st</sup> on TOP500 [14]. Core count trends for supercomputers are not necessarily bound to processor trends observed on Figure 1.1.

The problem of optimizing software for performance presents its own challenges. The search spaces that emerge from autotuning problems grow quickly to a size for which it would take a prohibitive amount of time to determine the best configuration by exhaustively evaluating all possibilities. Although this means we must seek to decrease the amount of possibilities, by restricting allowed parameter values, or dropping parameters completely, it is often unclear how to decide which parameters should be restricted or dropped. The next section introduces a simple autotuning problem, and presents an overview of autotuning search spaces from the size perspective. The section also briefly introduces methods used to explore autotuning search spaces, which are discussed in Chapter 3.

## 1.2 Characterizing Search Spaces from Autotuning Problems




Figure 1.4: Tiling and unrolling access patterns

```

int N = 256;
int BLOCK = 4;

int A[N][N], B[N][N], C[N][N];

int i, j, k;

for (i = 0; i < N; i++){
    for (j = 0; j < N; j++){
        C[i][j] = 0;

        for (k = 0; k < N; k++){
            C[i][j] += A[i][k] * B[k][j];
        }
    }
}

```

Listing 1.1: Test

## 1.3 Thesis Contributions

## 1.4 Text Structure

Chapter 3 presents background for the application of methods derived from search heuristics, mathematical optimization, and statistical learning to autotuning.

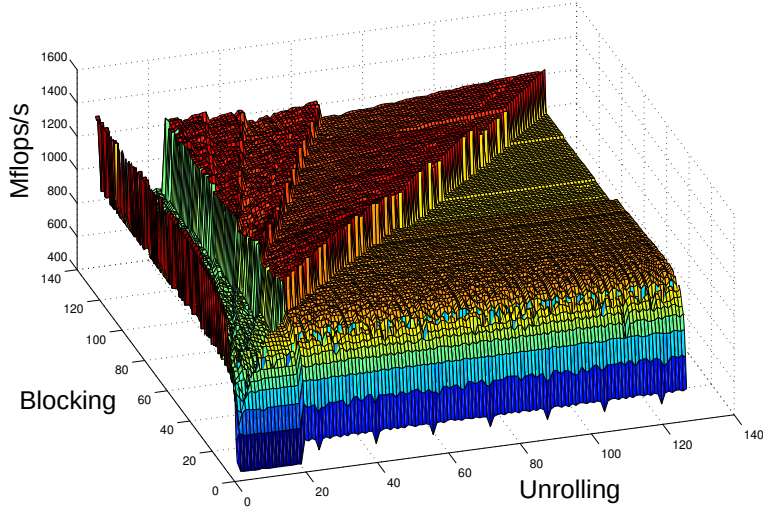


Figure 1.5: A search space defined for loop blocking and unrolling parameters, in a matrix multiplication kernel [16]




Figure 1.6: Dimension and search space size for autotuning problems from different domains [17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 16]. The left panel shows a zoomed view of the right panel



# **Reproducible Science**



# **Search and Statistical Learning for Autotuning**





# **Design of Experiments**



# Gaussian Process Regression

## 5.1 References



# Bibliography

- [1] D. C. Montgomery, *Design and analysis of experiments*. John Wiley & Sons, 2017.
- [2] C. K. Williams and C. E. Rasmussen, *Gaussian processes for machine learning*. MIT Press Cambridge, MA, 2006, vol. 2, no. 3.
- [3] P. E. Ceruzzi, E. Paul *et al.*, *A history of modern computing*. MIT Press, 2003.
- [4] O. R. N. Laboratory, "Summit – Oak Ridge Leadership Computing Facility," <https://www.olcf.ornl.gov/olcf-resources/compute-systems/summit>, 2020, [Online; accessed 08-Jun-2020].
- [5] G. E. Moore *et al.*, "Cramming more components onto integrated circuits," 1965.
- [6] Wikipedia, "Transistor count," [https://en.wikipedia.org/wiki/Transistor\\_count](https://en.wikipedia.org/wiki/Transistor_count), 2020, [Online; accessed 08-Jun-2020].
- [7] —, "Microprocessor Chronology," [https://en.wikipedia.org/wiki/Microprocessor\\_chronology](https://en.wikipedia.org/wiki/Microprocessor_chronology), 2020, [Online; accessed 08-Jun-2020].
- [8] R. H. Dennard, F. H. Gaensslen, H.-N. Yu, V. L. Rideout, E. Bassous, and A. R. LeBlanc, "Design of ion-implanted mosfet's with very small physical dimensions," *IEEE Journal of solid-state circuits*, vol. 9, no. 5, pp. 256–268, 1974.
- [9] D. J. Frank, R. H. Dennard, E. Nowak, P. M. Solomon, Y. Taur, and H.-S. P. Wong, "Device scaling limits of si mosfets and their application dependencies," *Proceedings of the IEEE*, vol. 89, no. 3, pp. 259–288, 2001.
- [10] H. Esmaeilzadeh, E. Blem, R. S. Amant, K. Sankaralingam, and D. Burger, "Dark silicon and the end of multicore scaling," in *2011 38th Annual international symposium on computer architecture (ISCA)*. IEEE, 2011, pp. 365–376.

- [11] S. Venkataramani, S. T. Chakradhar, K. Roy, and A. Raghunathan, "Approximate computing and the quest for computing efficiency," in *2015 52nd ACM/EDAC/IEEE Design Automation Conference (DAC)*. IEEE, 2015, pp. 1–6.
- [12] H.-Y. Cheng, J. Zhan, J. Zhao, Y. Xie, J. Sampson, and M. J. Irwin, "Core vs. uncore: The heart of darkness," in *2015 52nd ACM/EDAC/IEEE Design Automation Conference (DAC)*. IEEE, 2015, pp. 1–6.
- [13] J. Henkel, H. Khdr, S. Pagani, and M. Shafique, "New trends in dark silicon," in *2015 52nd ACM/EDAC/IEEE Design Automation Conference (Dac)*. IEEE, 2015, pp. 1–6.
- [14] Top500, "Top500 List," <https://www.top500.org/>, 2020, [Online; accessed 08-Jun-2020].
- [15] J. J. Dongarra, P. Luszczek, and A. Petitet, "The LINPACK benchmark: past, present and future," *Concurrency and Computation: practice and experience*, vol. 15, no. 9, pp. 803–820, 2003.
- [16] K. Seymour, H. You, and J. Dongarra, "A comparison of search heuristics for empirical code optimization." in *CLUSTER*, 2008, pp. 421–429.
- [17] P. Balaprakash, S. M. Wild, and B. Norris, "SPAPT: Search problems in automatic performance tuning," *Procedia Computer Science*, vol. 9, pp. 1959–1968, 2012.
- [18] J. Ansel, S. Kamil, K. Veeramachaneni, J. Ragan-Kelley, J. Bosboom, U.-M. O'Reilly, and S. Amarasinghe, "Opentuner: An extensible framework for program auto-tuning," in *Proceedings of the 23rd international conference on Parallel architectures and compilation*. ACM, 2014, pp. 303–316.
- [19] J.-H. Byun, R. Lin, K. A. Yelick, and J. Demmel, "Autotuning sparse matrix-vector multiplication for multicore," *EECS, UC Berkeley, Tech. Rep*, 2012.
- [20] F. Petrovič, D. Střelák, J. Hozzová, J. Ol'ha, R. Trembecký, S. Benkner, and J. Filipovič, "A benchmark set of highly-efficient cuda and opencl kernels and its dynamic autotuning with kernel tuning toolkit," *Future Generation Computer Systems*, 2020.
- [21] P. Balaprakash, M. Salim, T. Uram, V. Vishwanath, and S. Wild, "Deephyper: Asynchronous hyperparameter search for deep neural networks," in *2018 IEEE 25th*

- International Conference on High Performance Computing (HiPC)*. IEEE, 2018, pp. 42–51.
- [22] P. Bruel, S. Quinito Masnada, B. Videau, A. Legrand, J.-M. Vincent, and A. Goldman, “Autotuning under Tight Budget Constraints: A Transparent Design of Experiments Approach,” in *The 19th Annual IEEE/ACM International Symposium in Cluster, Cloud, and Grid Computing (CCGrid 2019)*. IEEE/ACM, 2019.
  - [23] P. Bruel, M. Amarís, and A. Goldman, “Autotuning gpu compiler parameters using opentuner,” in *Proceedings of the WSCAD (Simpósio em Sistemas Computacionais de Alto Desempenho)*, 2015.
  - [24] —, “Autotuning CUDA compiler parameters for heterogeneous applications using the OpenTuner framework,” *Concurrency and Computation: Practice and Experience*, pp. e3973–n/a, 2017.
  - [25] A. Mametjanov, P. Balaprakash, C. Choudary, P. D. Hovland, S. M. Wild, and G. Sabin, “Autotuning fpga design parameters for performance and power,” in *Field-Programmable Custom Computing Machines (FCCM), 2015 IEEE 23rd Annual International Symposium on*. IEEE, 2015, pp. 84–91.
  - [26] C. Xu, G. Liu, R. Zhao, S. Yang, G. Luo, and Z. Zhang, “A parallel bandit-based approach for autotuning fpga compilation,” in *Proceedings of the 2017 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays*. ACM, 2017, pp. 157–166.
  - [27] A. Tiwari, C. Chen, J. Chame, M. Hall, and J. K. Hollingsworth, “A scalable autotuning framework for compiler optimization,” in *Parallel & Distributed Processing, 2009. IPDPS 2009. IEEE International Symposium on*. IEEE, 2009, pp. 1–12.
  - [28] F. Hutter, H. H. Hoos, K. Leyton-Brown, and T. Stützle, “ParamILS: an automatic algorithm configuration framework,” *Journal of Artificial Intelligence Research*, vol. 36, no. 1, pp. 267–306, 2009.
  - [29] Y. Chu, C. Luo, H. H. Hoos, Q. Lin, and H. You, “Improving the performance of stochastic local search for maximum vertex weight clique problem using programming by optimization,” *arXiv*, pp. arXiv–2002, 2020.
  - [30] I. Tuzov, D. de Andrés, and J.-C. Ruiz, “Tuning synthesis flags to optimize implementation goals: Performance and robustness of the leon3 processor as a case study,” *Journal of Parallel and Distributed Computing*, vol. 112, pp. 84–96, 2018.

- [31] M. M. Ziegler, H.-Y. Liu, G. Gristede, B. Owens, R. Nigaglioni, J. Kwon, and L. P. Carloni, "Syntunsys: A synthesis parameter autotuning system for optimizing high-performance processors," in *Machine Learning in VLSI Computer-Aided Design*. Springer, 2019, pp. 539–570.
- [32] M. Gerndt, S. Benkner, E. César, C. Navarrete, E. Bajrovic, J. Dokulil, C. Guillén, R. Mijakovic, and A. Sikora, "A multi-aspect online tuning framework for hpc applications," *Software Quality Journal*, vol. 26, no. 3, pp. 1063–1096, 2018.
- [33] J. Kwon, M. M. Ziegler, and L. P. Carloni, "A learning-based recommender system for autotuning design flows of industrial high-performance processors," in *2019 56th ACM/IEEE Design Automation Conference (DAC)*. IEEE, 2019, pp. 1–6.
- [34] T. Wang, N. Jain, D. Beckingsale, D. Boehme, F. Mueller, and T. Gamblin, "Funcy-tuner: Auto-tuning scientific applications with per-loop compilation," in *Proceedings of the 48th International Conference on Parallel Processing*, 2019, pp. 1–10.
- [35] J. Ol'ha, J. Hozzová, J. Fousek, and J. Filipovič, "Exploiting historical data: pruning autotuning spaces and estimating the number of tuning steps," in *European Conference on Parallel Processing*. Springer, 2019, pp. 295–307.